

# **Digitale Bildverarbeitung und Mustererkennung**

Deep Learning Foundations

**Classification & Object Detection and  
Transfer Learning**

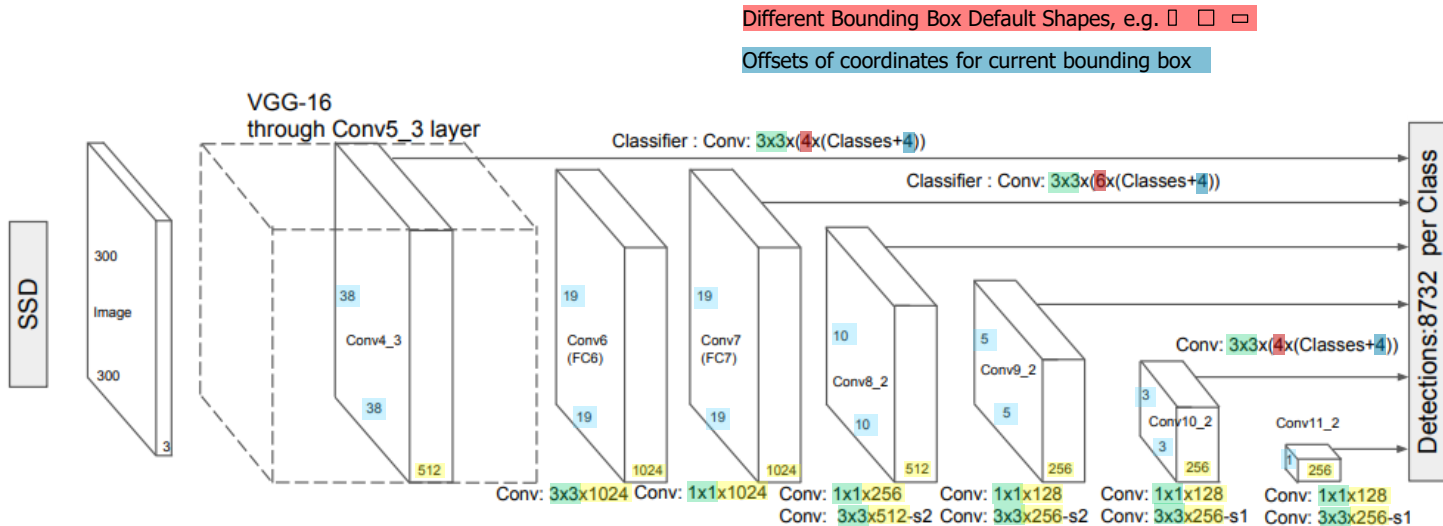
Segmentation Networks

Deep Reinforcement Learning

Generative Adversarial Networks

Recurrent Neural Networks

## Neural Network Object Detection – Single Shot Detector with VGG-16 Backbone



## Kernel Size

$$\text{Amount of Kernels in Layer} = \frac{\text{Depth of subsequent Feature Map}}{\text{Width \& Height of Feature Map}}$$

Liu, W., Anguelov, D., Erhan, D., Szegedy, C., Reed, S., Fu, C. Y., & Berg, A. C. (2016, October). SSD: Single shot multibox detector. In *European conference on computer vision* (pp. 21-37). Springer, Cham. (<https://arxiv.org/pdf/1512.02325.pdf>)

# Neural Network Object Detection - Basic structure

## SSD: Single Shot MultiBox Detector

### SSD: Single Shot MultiBox Detector

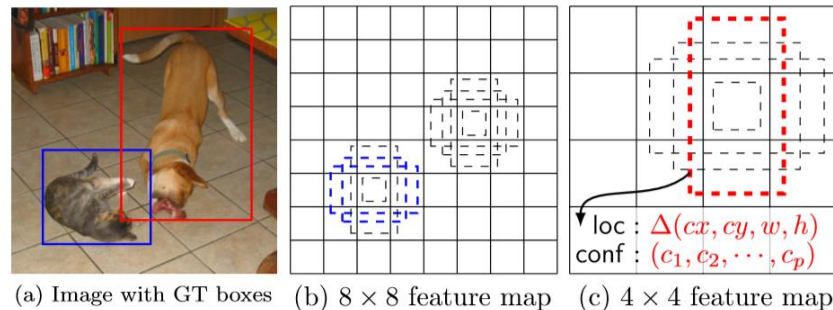
Wei Liu<sup>1</sup>, Dragomir Anguelov<sup>2</sup>, Dumitru Erhan<sup>3</sup>, Christian Szegedy<sup>3</sup>,  
Scott Reed<sup>4</sup>, Cheng-Yang Fu<sup>1</sup>, Alexander C. Berg<sup>1</sup>

<sup>1</sup>UNC Chapel Hill <sup>2</sup>Zoox Inc. <sup>3</sup>Google Inc. <sup>4</sup>University of Michigan, Ann-Arbor  
wliu@cs.unc.edu, <sup>2</sup>drago@zoox.com, <sup>3</sup>{dumitru, szegedy}@google.com,  
<sup>4</sup>reedscot@umich.edu, <sup>1</sup>{cyfu, aberg}@cs.unc.edu

Discretization of the input image into a SxS grid of different sizes.

Fully Convolutional Network predicts class scores and box offsets for given default bounding boxes per size.

Final detections are chosen with non-maximum suppression.



<https://arxiv.org/abs/1512.02325>

# Neural Network Object Detection - Basic structure

## RFCN: Region-based Fully Convolutional Network

Two stage object detection.

First stage is Fully Convolutional Network, such as ResNet-101 for region proposals

Second stage does classification on the max-pooled proposed regions

---

## R-FCN: Object Detection via Region-based Fully Convolutional Networks

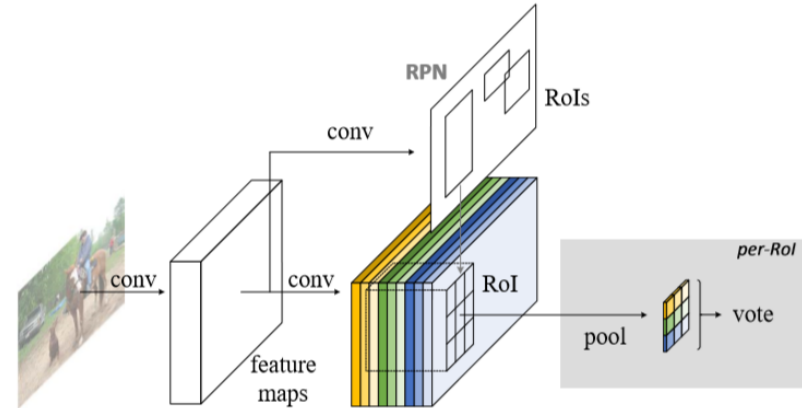
---

Jifeng Dai  
Microsoft Research

Yi Li\*  
Tsinghua University

Kaiming He  
Microsoft Research

Jian Sun  
Microsoft Research



<https://arxiv.org/abs/1605.06409>

# Object Detection with neural networks

## TensorFlow Model Zoo

Collection of detection models pre-trained on different object detection datasets.



| Model name                                    | Speed (ms) | COCO mAP  | Outputs     |
|---|------------|-----------|-------------|
| SSD MobileNet v2 320x320                      | 19         | 20.2      | Boxes       |
| SSD MobileNet V1 FPN 640x640                  | 48         | 29.1      | Boxes       |
| SSD MobileNet V2 FPNLite 320x320              | 22         | 22.2      | Boxes       |
| SSD MobileNet V2 FPNLite 640x640              | 39         | 28.2      | Boxes       |
| SSD ResNet50 V1 FPN 640x640 (RetinaNet50)     | 46         | 34.3      | Boxes       |
| SSD ResNet50 V1 FPN 1024x1024 (RetinaNet50)   | 87         | 38.3      | Boxes       |
| SSD ResNet101 V1 FPN 640x640 (RetinaNet101)   | 57         | 35.6      | Boxes       |
| SSD ResNet101 V1 FPN 1024x1024 (RetinaNet101) | 104        | 39.5      | Boxes       |
| SSD ResNet152 V1 FPN 640x640 (RetinaNet152)   | 80         | 35.4      | Boxes       |
| SSD ResNet152 V1 FPN 1024x1024 (RetinaNet152) | 111        | 39.6      | Boxes       |
| Faster R-CNN ResNet50 V1 640x640              | 53         | 29.3      | Boxes       |
| Faster R-CNN ResNet50 V1 1024x1024            | 65         | 31.0      | Boxes       |
| Faster R-CNN ResNet50 V1 800x1333             | 65         | 31.6      | Boxes       |
| Faster R-CNN ResNet101 V1 640x640             | 55         | 31.8      | Boxes       |
| Faster R-CNN ResNet101 V1 1024x1024           | 72         | 37.1      | Boxes       |
| Faster R-CNN ResNet101 V1 800x1333            | 77         | 36.6      | Boxes       |
| Faster R-CNN ResNet152 V1 640x640             | 64         | 32.4      | Boxes       |
| Faster R-CNN ResNet152 V1 1024x1024           | 85         | 37.6      | Boxes       |
| Faster R-CNN ResNet152 V1 800x1333            | 101        | 37.4      | Boxes       |
| Faster R-CNN Inception ResNet V2 640x640      | 206        | 37.7      | Boxes       |
| Faster R-CNN Inception ResNet V2 1024x1024    | 236        | 38.7      | Boxes       |
| Mask R-CNN Inception ResNet V2 1024x1024      | 301        | 39.0/34.6 | Boxes/Masks |

[https://github.com/tensorflow/models/blob/master/research/object\\_detection/g3doc/tf2\\_detection\\_zoo.md](https://github.com/tensorflow/models/blob/master/research/object_detection/g3doc/tf2_detection_zoo.md)

# References

- [2] Joseph Redmon, Santosh Kumar Divvala, Ross B. Girshick, and Ali Farhadi. You only look once: Unified, real-time object detection. *CoRR*, abs/1506.02640, 2015.
- [3] Wei Liu, Dragomir Anguelov, Dumitru Erhan, Christian Szegedy, Scott Reed, Cheng-Yang Fu, and Alexander C. Berg. Ssd: Single shot multibox detector. *14th european conference on computer vision*, pages 21–37, 2016.
- [4] Jifeng Dai, Yi Li, Kaiming He, and Jian Sun. R-FCN: object detection via region-based fully convolutional networks. *CoRR*, abs/1605.06409, 2016.
- [5] Tsung-Yi Lin, Michael Maire, Serge J. Belongie, Lubomir D. Bourdev, Ross B. Girshick, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C. Lawrence Zitnick. Microsoft COCO: common objects in context. *CoRR*, abs/1405.0312, 2014.
- [6] Mark Everingham, Luc Van Gool, Christopher K. I. Williams, John Winn, and Andrew Zisserman. The pascal visual object classes (voc) challenge. *International Journal of Computer Vision*, 88(2):303–338, Jun 2010.
- [7] Andreas Geiger, Philip Lenz, and Raquel Urtasun. Are we ready for autonomous driving? the kitti vision benchmark suite. In *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2012.
- [8] Jesse Davis and Mark Goadrich. The relationship between precision-recall and roc curves. In *Proceedings of the 23rd International Conference on Machine Learning, ICML '06*, pages 233–240, New York, NY, USA, 2006. ACM.
- [9] Mark Everingham, Luc Van Gool, Christopher K. I. Williams, John Winn, and Andrew Zisserman. The pascal visual object classes (voc) challenge. *International Journal of Computer Vision*, 88(2):303–338, Jun 2010.

This lecture in one slide

Object Detection with neural networks

**Transfer Learning with neural networks**

Fine Tuning

Freezing Weights



## Transfer Learning with Neural Networks

The intuition behind transfer learning is that if a model trained on a large and general enough dataset, this model will effectively serve as a **generic model of the visual world**.

## Transfer Learning with Neural Networks

The intuition behind transfer learning is that if a model trained on a large and general enough dataset, this model will effectively serve as a generic model of the visual world.

*You can then take advantage of these learned features without having to start from scratch training a large model on a large dataset.*

## Transfer Learning with Neural Networks - Fine Tuning

A **Pre-trained model** is a model that is trained on the source domain for a source task.



Pre-trained  
model

Car

## Transfer Learning with Neural Networks - Fine Tuning

A **Pre-trained model** is a model that is trained on source domain for a source task.

Think **Object Detection Model Zoo**



Pre-trained  
model

Car

[https://github.com/tensorflow/models/blob/master/research/object\\_detection/g3doc/tf2\\_detection\\_zoo.md](https://github.com/tensorflow/models/blob/master/research/object_detection/g3doc/tf2_detection_zoo.md)

## Transfer Learning with Neural Networks - Fine Tuning

A **Pre-trained model** is a model that is trained on the source domain for a source task.

The source domain is usually a **large dataset**.



Pre-trained  
model

Car

## Transfer Learning with Neural Networks - Fine Tuning

A **Pre-trained model** is a model that is trained on the source domain for a source task.

The source domain is usually a **large dataset**.

Think of **COCO**, **ImageNet**, **PASCAL VOC**



Pre-trained  
model

Car

## Transfer Learning with Neural Networks - Fine Tuning

You either use a **pre-trained model as it is**.



Pre-trained  
model

Car

## Transfer Learning with Neural Networks - Fine Tuning

You either use a pre-trained model as it is.

- No additional training
- Usually **performance average or low**



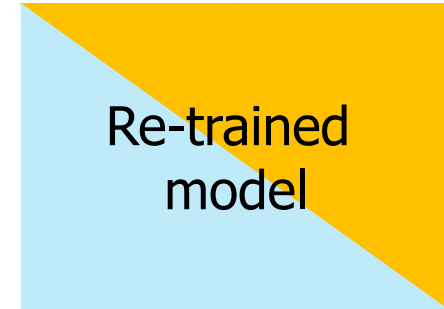
Pre-trained  
model

Car



## Transfer Learning with Neural Networks - Fine Tuning

You can **fine-tune** a **pre-trained model** by retraining on additional data.

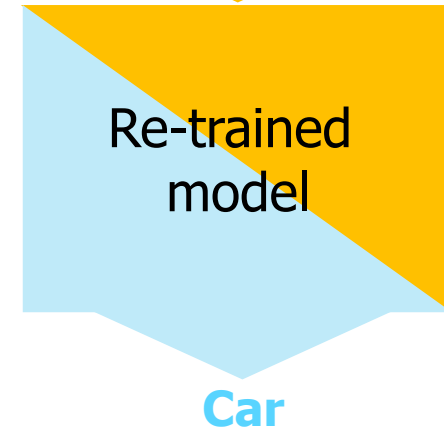


Car

## Transfer Learning with Neural Networks - Fine Tuning

You can fine-tune a pre-trained model by retraining on additional data.

- Additional work to set up retraining pipeline
- Usually performance drops in the source domain
- **Performance improvement** in the target domain



## Transfer Learning with Neural Networks - Feature Extraction

You can **repurpose** the **pre-trained layers** as feature extraction layers for low level features.

And adding layers that learn the required high level features on the new data.



Pre-trained  
layers

Newly-trained  
layers

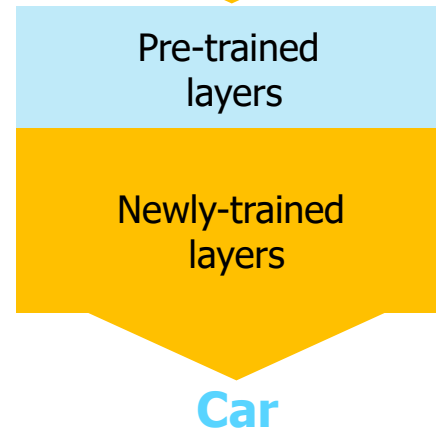
Car

## Transfer Learning with Neural Networks - Feature Extraction

You can repurpose the pre-trained layers as feature extraction layers for low level features.

And adding layers that learn the required high level features on the new data.

- Additional work to assemble layer structure
- **Regularization** effects
- **Performance improvement** in the target domain
- **Output** layer becomes **adjustable** (such as for adding a class)

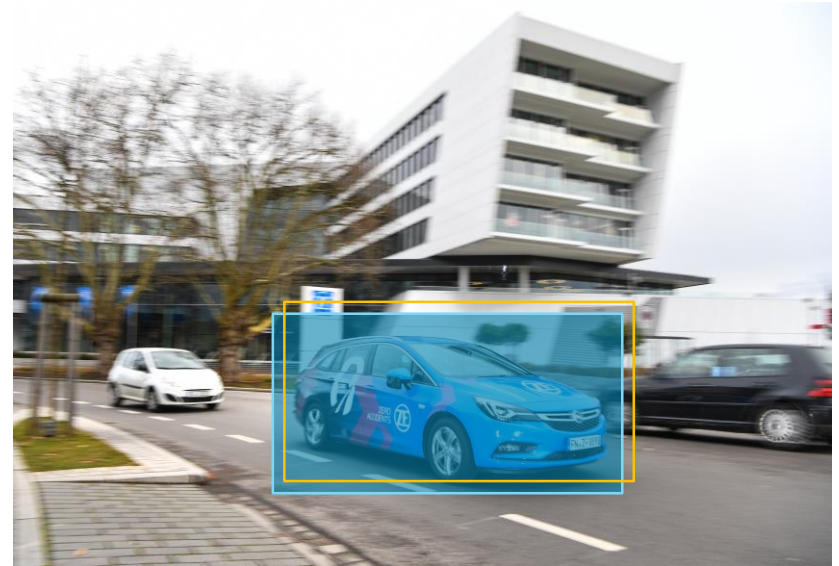


# Neural Network Object Detection - mAP

## So how do we evaluate correctness of a detection?

Define when a detection is correct and when it is not.

A prediction will usually not overlap perfectly with the ground truth.

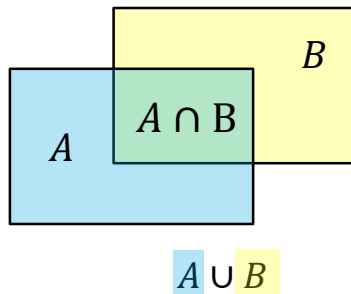


# Neural Network Object Detection - mAP

## So how do we evaluate correctness of a detection?

Define when a detection is correct, a true positive and when it is not.  
A prediction will usually not overlap perfectly with the ground truth.

A correct detection is based on  
**Intersection over Union (IoU)** with the  
ground truth



$$IoU(A, B) = \frac{A \cap B}{A \cup B}$$

$$IoU \geq p, \quad p \in (0,1)$$

# Neural Network Object Detection - mAP

## Precision

$$P = \frac{TP}{TP + FP}$$

*TP = True Positives*

*FP = False Positives*

*P = Precision*

Correct detections over the number of predicted detections.

## Recall

$$R = \frac{TP}{TP + FN}$$

*FN = False Negatives*

*R = Recall*

Correct detections over the number of groundtruth detection.

## F1 Score

$$F1 = 2 \cdot \frac{P \cdot R}{P + R}$$

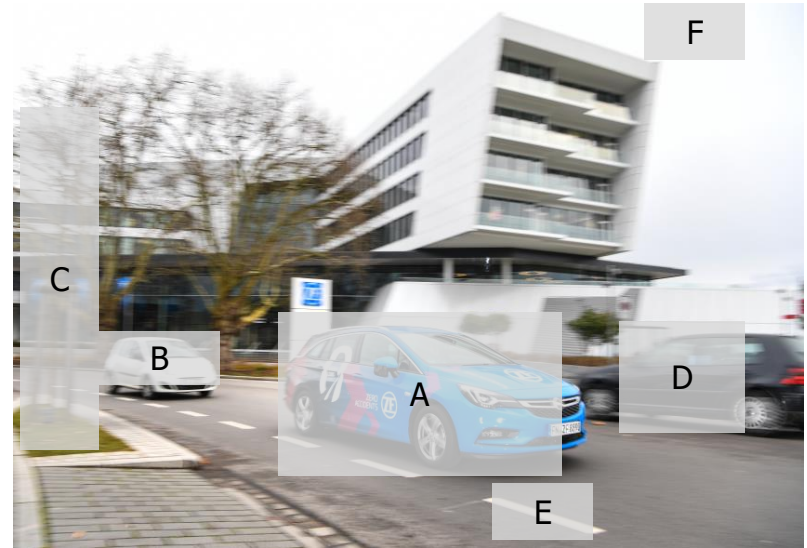
Harmonic mean of precision and recall.

# Neural Network Object Detection - mAP

## Precision Recall Curve

By moving the probability threshold of the outputs of the model we can fill in a precision recall curve. When deploying a model you would need to set exactly one probability threshold.

| Box | Prob. | Correct | TP | FP | FN | Prec. | Rec. |
|-----|-------|---------|----|----|----|-------|------|
| A   | 0.95  | True    | 1  | 0  | 2  | 1     | 0.33 |
| B   | 0.85  | True    | 2  | 0  | 1  | 1     | 0.66 |
| C   | 0.7   | False   | 2  | 1  | 1  | 0.66  | 0.66 |
| D   | 0.5   | True    | 3  | 1  | 0  | 0.75  | 1    |
| E   | 0.3   | False   | 3  | 2  | 0  | 0.6   | 1    |
| F   | 0.2   | False   | 3  | 3  | 0  | 0.5   | 1    |





# Neural Network Object Detection - mAP

## Precision Recall Curve

By moving the probability threshold of the outputs of the model we can fill in a precision recall curve. When deploying a model you would need to set exactly one probability threshold.

## Average Precision (AP)

AP is the area under the curve of the Precision Recall Curve.

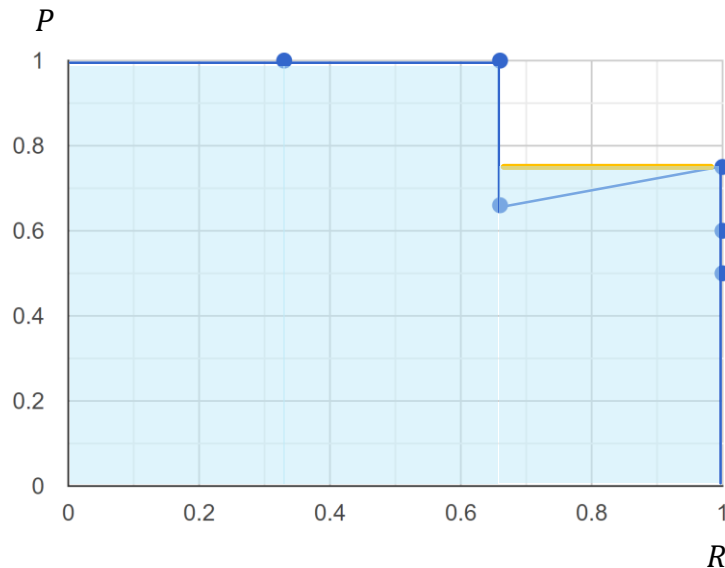
$$AP = \sum_n^N (R_n - R_{n-1})P_n$$

$R_n$  = Recall at threshold step  $n$

$P_n$  = Precision at threshold step  $n$

| Prec. | Rec. |
|-------|------|
| 1     | 0.33 |
| 1     | 0.66 |
| 0.66  | 0.66 |
| 0.75  | 1    |
| 0.6   | 1    |
| 0.5   | 1    |

$$AP = 0.33 + 0.33 + 0 + 0.26 + 0 + 0 = 0.92$$



# Neural Network Object Detection - mAP

## Precision Recall Curve

By moving the probability threshold of the outputs of the model we can fill in precision recall curve.  
When deploying a model you would need to set exactly one threshold.

## Average Precision (AP)

AP is the area under the curve of the Precision Recall Curve.

$$AP = \sum_n^N (R_n - R_{n-1})P_n$$

$R_n$  = Recall at threshold step  $n$

$P_n$  = Precision at threshold step  $n$

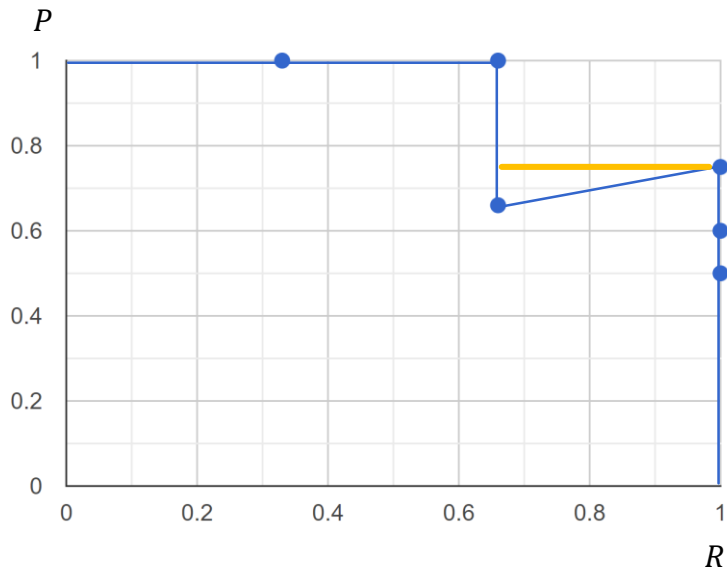
## mean Average Precision (mAP)

Average Precision over different classes.

$$mAP = \frac{1}{N} \sum_c^N AP_c$$

$c$  = current class

$N$  = amount of classes

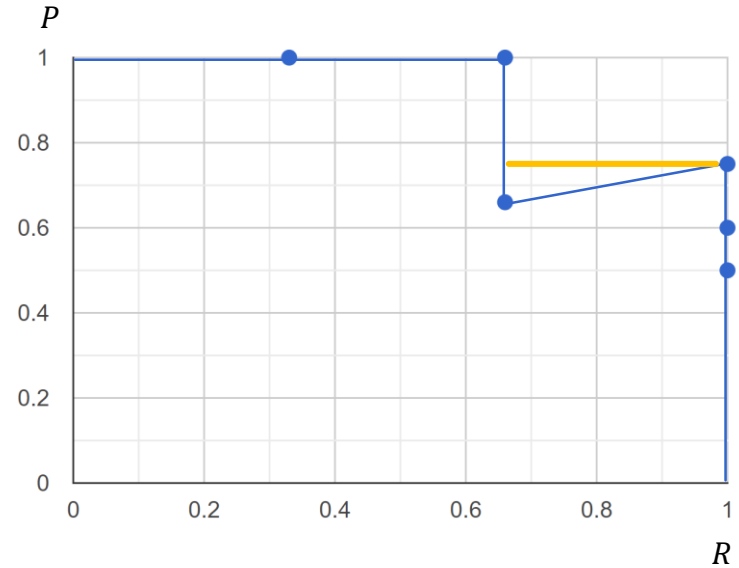


# Neural Network Object Detection – COCO mAP

**mAP@[0.5:0.05:0.95]**

COCO mAP.

Average of Mean Average Precision over a set of IoU levels (0.5, 0.55, 0.6,..., 0.95).



Deep Learning Foundations

Classification & Object Detection and Transfer Learning

**Segmentation Networks**

Deep Reinforcement Learning

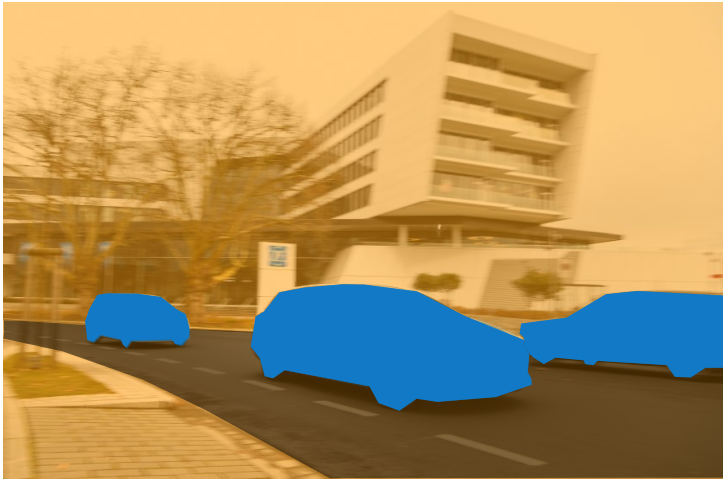
Generative Adversarial Networks

Recurrent Neural Networks

## Introduction – Segmentation a problem statement

*Segmentation is classification on pixel-level, which results in super-pixels or segments or groups of pixels based on some criteria.*

### Semantic Segmentation



### Instance Segmentation



### Autonomous Driving

- Scene understanding
- Understanding of shapes
- Free space detection

Geo Analytics

Medical Imaging



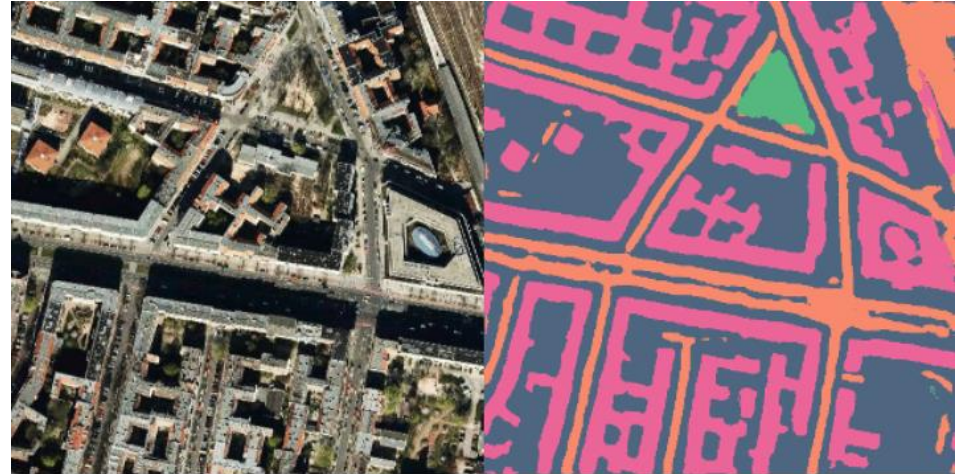
<https://www.cityscapes-dataset.com/>

# Introduction – Applications for segmentation

## Autonomous Driving

### Geo Analytics

- Building structures
- Road network analysis
- Wildfire detection
- Water supply tracking
- Real time crisis management
- Weather prediction



<https://github.com/mapbox/robosat>

## Medical Imaging

# Introduction – Applications for segmentation

Autonomous Driving

Geo Analytics

## Medical Imaging

- Tissue localization and analysis
- Volume approximations
- Surgery planning
- Temporal tumor or tissue development
- Tooling for drug testing



<https://osf.io/snb6p/>



## Introduction – Conventional segmentation approaches

In order to design neural networks it is a good thing to really understand the task at hand.

Thresholding

Edge detection

Clustering

## Introduction – Conventional segmentation approaches

### Thresholding

The simplest method of image segmentation is thresholding.

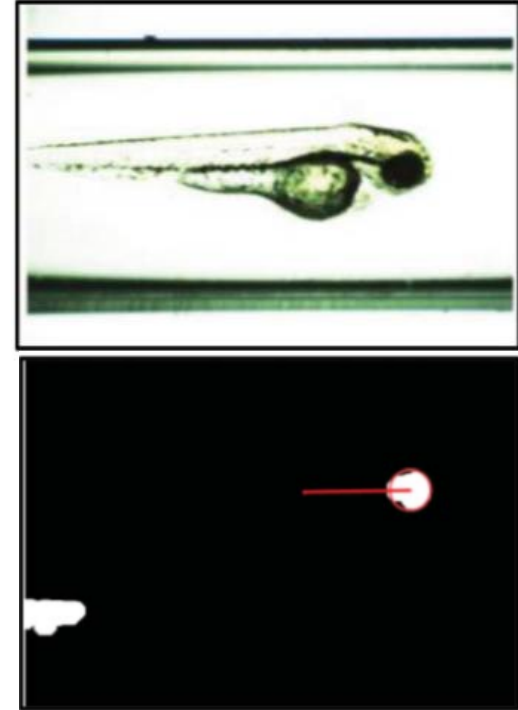
This method is based on a threshold value to turn a gray-scale image into a binary image (mask).

Usually this is just one step of many.

Edge detection

Clustering

Region growing



*Thresholding on Zebrafish for eye segmentation*

# Lena Test image

Lena, the 'hello world!' of image processing. 330x330

Cover photo of 1972 Playboy magazine of the Swedish model Lena Söderberg.

Since then Lena was a guest at several IEEE conferences. The image also sparked discussions on gender-equality in the male-dominated field of engineering.

It is a good test image because of its detail, flat regions, shading, and texture.



<https://en.wikipedia.org/wiki/Lenna>

# Introduction – Conventional segmentation approaches

Thresholding

## Edge detection

Segment boundaries and edges are closely related.

Since there is often a large gradient at the segment boundaries.



*Canny Edge Detection (multi-step approach)*

$$L_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} L \quad \text{and} \quad L_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} L.$$

*Sobel Operators for Edge Detection*

Clustering

## Introduction – Conventional segmentation approaches

Thresholding

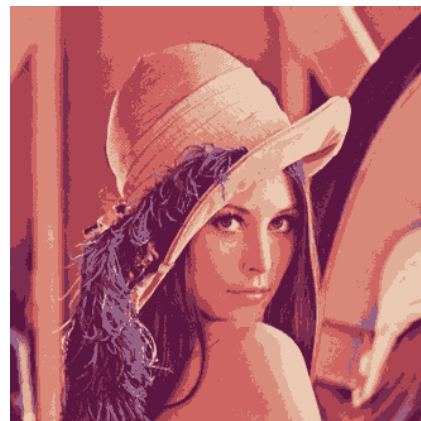
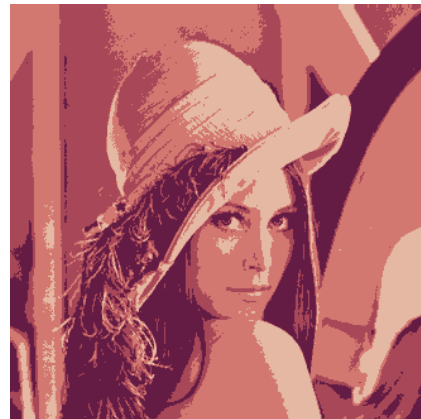
Edge detection

### Clustering (Color quantization)

K-means with 3 features (R,G,B) and K centroids.

The centroids are iteratively adjusted until convergence.

After the clustering, the centroid values are applied to the pixels in their cluster.



*Clustering for K=4 (top) and K=8 (bottom)*

# References

- [1] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [2] Andrej Karpathy. Cs231n: Convolutional neural networks for visual recognition. <http://cs231n.github.io/neural-networks-3/>, 2018. Zugriff: 20.01.2018.
- [3] Schutera, Mark, Steffen Just, Jakob Gierden, Ralf Mikut, Markus Reischl, and Christian Pylatiuk. 2019. "Machine Learning Methods for Automated Quantification of Ventricular Dimensions." OSF. March 28. [osf.io/snb6p](https://osf.io/snb6p).
- [4] Mark Schutera, Thomas Dickmeis, Marina Mione, Ravindra Peravali, Daniel Marcato, Markus Reischl, Ralf Mikut, and Christian Pylatiuk. Automated phenotype pattern recognition of zebrafish for high-throughput screening. *Bioengineered*, 7(4):261–265, 2016.
- [5] Canny, J., *A Computational Approach To Edge Detection*, IEEE Transactions on Pattern Analysis and Machine Intelligence, 8(6):679–698, 1986.

## Introduction to segmentation

### **Segmentation with neural networks**

Basic structure

Overview state-of-the-art

Datasets and benchmarking

Deep dive U-Net

# Neural Network Segmentation - Basic structure

## Why classical segmentation approaches

- Interpretability
- Only a few samples needed
- No labeling needed
- No training needed
- Usually better runtime during inference

Why not?



# Neural Network Segmentation - Basic structure

Why classical segmentation approaches

## **Why segmentation by neural networks?**

- Do generalize better
- Feature engineering has a limited capacity to capture semantics
- Feature engineering is expensive and time consuming

# Neural Network Segmentation - Basic structure

## Feature Representation by Convolution

*Idea is to classify each pixel of an input image by representation learning*

Downsampling

Upsampling

Parameter sharing

### Feature Representation by Convolution

|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 2 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 2 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Zero-Padding is adding zero-valued pixel to the image border (gray area).

Downsampling

Upsampling

Parameter sharing

## Neural Network Segmentation - Basic structure

### Feature Representation by Convolution

|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 2 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 2 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Zero-padded image

|    |    |    |
|----|----|----|
| 0  | 0  | -1 |
| -1 | 0  | 0  |
| -1 | -1 | -1 |

Filter

|   |
|---|
| 0 |
|---|

Bias

|    |    |    |
|----|----|----|
| -4 | -4 | 0  |
| -3 | -4 | -3 |
| 0  | -3 | -1 |

Output

Downsampling

Upsampling

Parameter sharing

## Neural Network Segmentation - Basic structure

### Feature Representation by Convolution

|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 2 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 2 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

|    |    |    |
|----|----|----|
| 0  | 0  | -1 |
| -1 | 0  | 0  |
| -1 | -1 | -1 |

Filter

|   |
|---|
| 0 |
|---|

Bias

|    |    |    |
|----|----|----|
| -4 | -4 | 0  |
| -3 | -4 | -3 |
| 0  | -3 | -1 |

Output

Amount of filters or convolution depth: 1

Filter step size or Stride: 2

Zero-padded image

Downsampling

Upsampling

Parameter sharing

## Neural Network Segmentation - Basic structure

### Feature Representation by Convolution

|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 2 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 2 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

|    |    |    |
|----|----|----|
| 0  | 0  | -1 |
| -1 | 0  | 0  |
| -1 | -1 | -1 |

Filter

|   |
|---|
| 0 |
|---|

Bias

|    |    |    |
|----|----|----|
| -4 | -4 | 0  |
| -3 | -4 | -3 |
| 0  | -3 | -1 |

Output

Downsampling

Upsampling

Parameter sharing

### Review edge detector:

Similar idea, now the parameters of the filters are learned. We want a lot of filters!

$$L_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} L \quad \text{and} \quad L_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} L.$$

## Neural Network Segmentation - Basic structure

### Feature Representation by Convolution

|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 1 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 0 | 0 | 0 | 0 |
| 0 | 2 | 2 | 2 | 1 | 0 | 0 |
| 0 | 1 | 0 | 0 | 2 | 0 | 0 |
| 0 | 1 | 1 | 0 | 1 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Downsampling

Upsampling

Parameter sharing

|    |    |    |
|----|----|----|
| 0  | 0  | -1 |
| -1 | 0  | 0  |
| -1 | -1 | -1 |

Filter

|   |
|---|
| 0 |
|---|

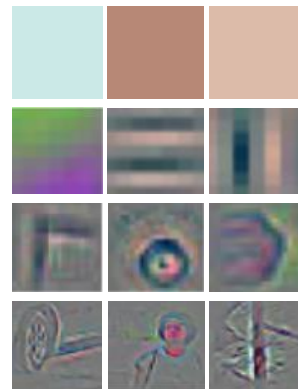
Bias

|    |    |    |
|----|----|----|
| -4 | -4 | 0  |
| -3 | -4 | -3 |
| 0  | -3 | -1 |

Output

#### Review edge detector:

Similar idea, now the parameters of the filters are learned. And we want to go deep!



# Neural Network Segmentation - Basic structure

Convolutions

## Downsampling

*Convolutions at original image resolution are computational expensive:*

*Filter dimensions  $\times$  image dimensions  $\times$  number of filters  $\times$  number of input channels.*

Motivating a convolutional  
encoder-decoder  
structure and Downsampling.

Upsampling

Parameter sharing



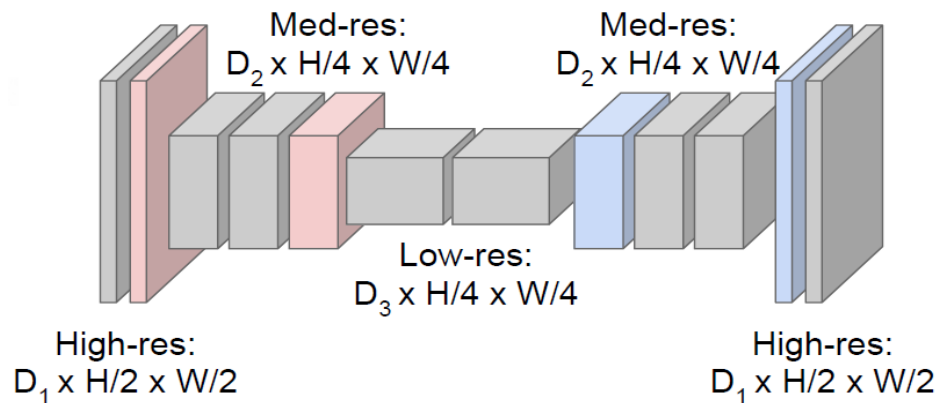
# Neural Network Segmentation - Basic structure

## Convolutions

### Downsampling

*Convolutions at original image resolution are computational expensive:  
Filter dimensions  $\times$  image dimensions  $\times$  number of filters  $\times$  number of input channels.*

Motivating a convolutional  
encoder-decoder  
structure and Downsampling.



<http://cs231n.github.io/>

## Upsampling

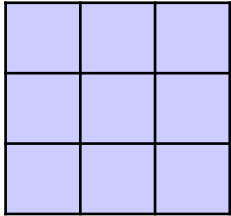
## Parameter sharing

# Neural Network Segmentation - Basic structure

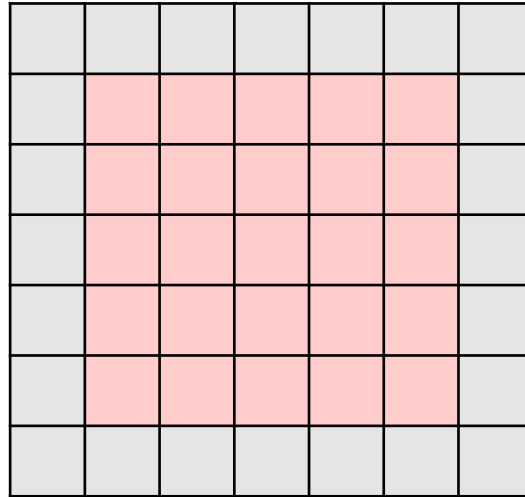
Convolutions

**Downsampling**

**Strided convolutions**



Filter 3x3x1



Zero-padded image

Upsampling

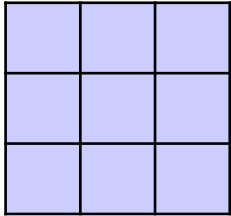
Parameter sharing

# Neural Network Segmentation - Basic structure

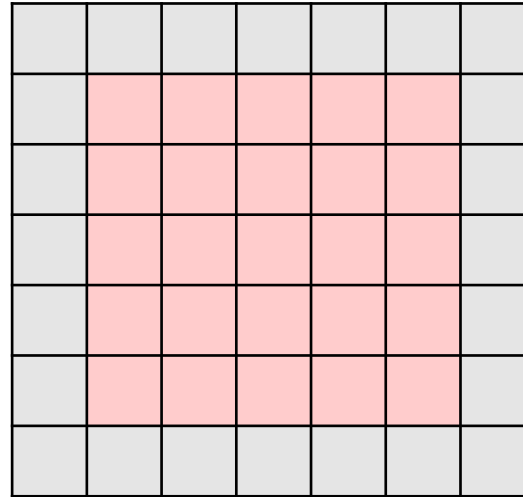
Convolutions

**Downsampling**

**Strided convolutions**

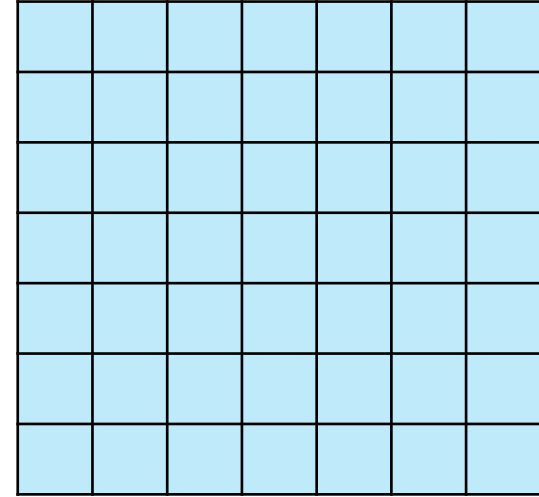


Filter 3x3x1



Zero-padded image

**Stride 1**



Upsampling

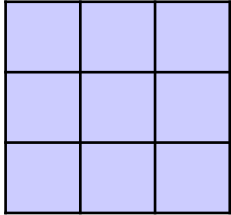
Parameter sharing

# Neural Network Segmentation - Basic structure

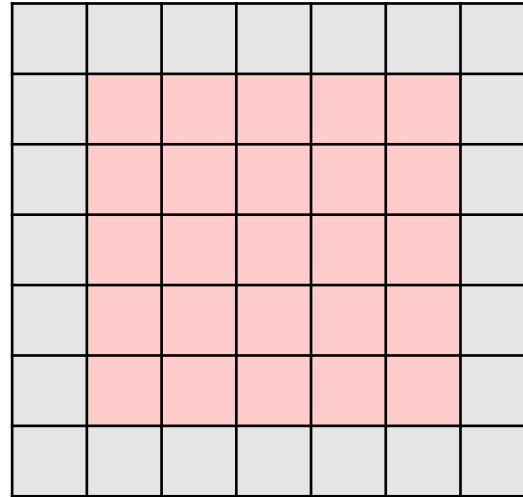
Convolutions

**Downsampling**

**Strided convolutions**

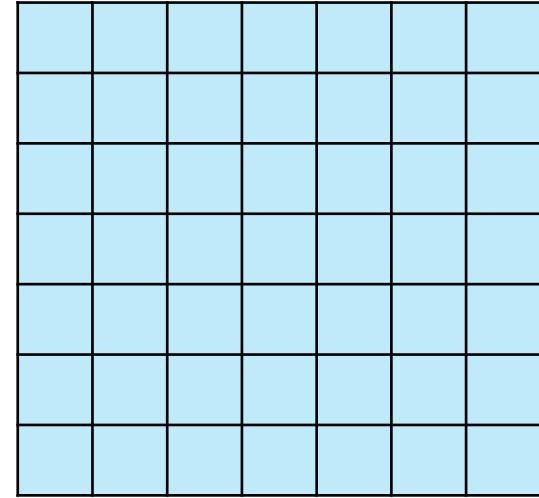


Filter 3x3x1

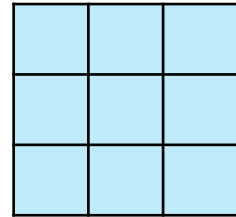


Zero-padded image

**Stride 1**



**Stride 2**



Upsampling

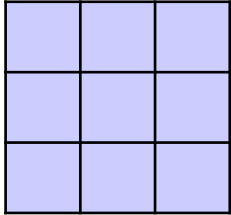
Parameter sharing

# Neural Network Segmentation - Basic structure

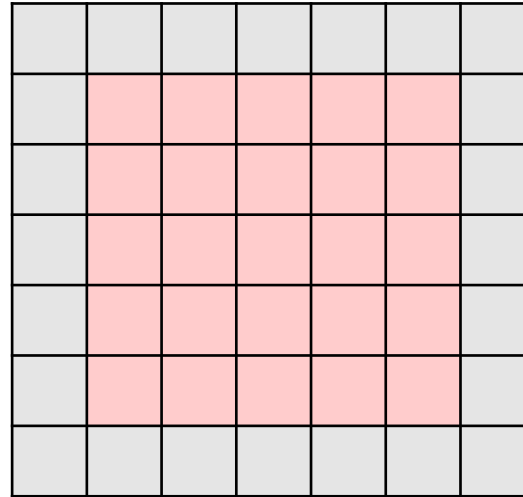
Convolutions

**Downsampling**

**Strided convolutions**



Filter 3x3x1

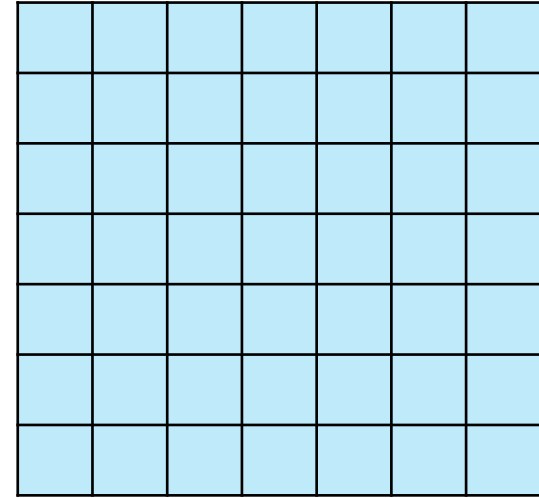


Zero-padded image

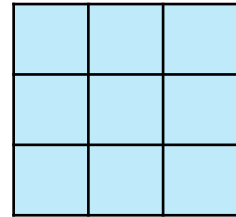
Upsampling

Parameter sharing

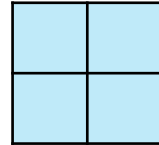
**Stride 1**



**Stride 2**



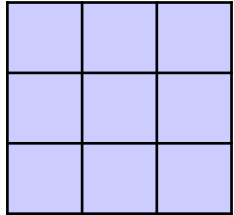
**Stride 4**



# Neural Network Segmentation - Basic structure

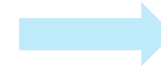
Convolutions

## Downsampling Max Pooling



Max Pooling  
3x3 Stride 3

|   |   |   |   |   |   |
|---|---|---|---|---|---|
| 5 | 4 | 6 | 3 | 1 | 7 |
| 3 | 4 | 1 | 6 | 4 | 4 |
| 6 | 2 | 5 | 6 | 6 | 4 |
| 2 | 6 | 9 | 8 | 6 | 3 |
| 4 | 8 | 5 | 6 | 8 | 2 |
| 3 | 1 | 7 | 8 | 6 | 3 |



|   |   |
|---|---|
| 6 | 7 |
| 9 | 8 |

Upsampling

Parameter sharing

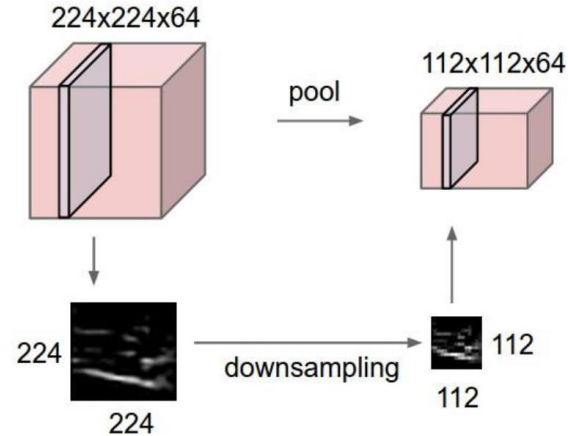
# Neural Network Segmentation - Basic structure

## Convolutions

### Downsampling

#### Max Pooling

Intuition is to decrease the resolution while keeping the strongest features of each channel.



<https://selfdrivingcars.mit.edu/>

## Upsampling

## Parameter sharing

# Neural Network Segmentation - Basic structure

Convolutions

Downsampling

## Upsampling

*Classification needs to happen in original image resolution*

Motivating Upsampling inside the network structure.

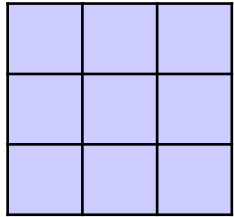
Parameter sharing



# Neural Network Segmentation - Basic structure

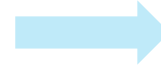
Convolutions  
Downsampling

**Upsampling**  
**Nearest neighbor**



3x3 Stride 3

|   |   |
|---|---|
| 6 | 7 |
| 9 | 8 |



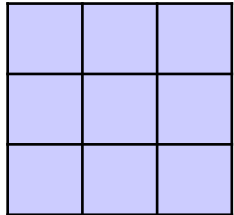
|   |   |   |   |   |   |
|---|---|---|---|---|---|
| 6 | 6 | 6 | 7 | 7 | 7 |
| 6 | 6 | 6 | 7 | 7 | 7 |
| 6 | 6 | 6 | 7 | 7 | 7 |
| 9 | 9 | 9 | 8 | 8 | 8 |
| 9 | 9 | 9 | 8 | 8 | 8 |
| 9 | 9 | 9 | 8 | 8 | 8 |

Parameter sharing

# Neural Network Segmentation - Basic structure

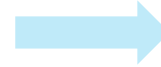
Convolutions  
Downsampling

## Upsampling Bed of Nails



3x3 Stride 3

|   |   |
|---|---|
| 6 | 7 |
| 9 | 8 |



|   |   |   |   |   |   |
|---|---|---|---|---|---|
| 6 | 0 | 0 | 7 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 |
| 9 | 0 | 0 | 8 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 |

Parameter sharing

# Neural Network Segmentation - Basic structure

Convolutions

Downsampling

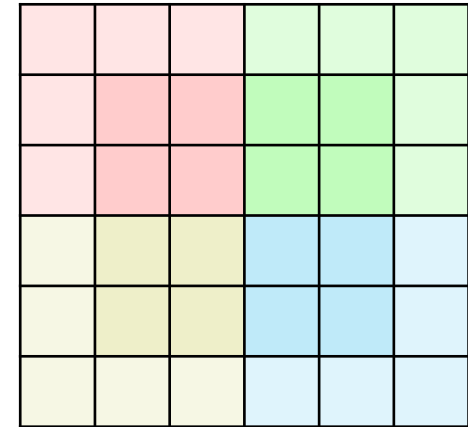
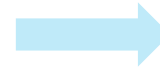
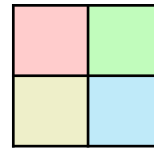
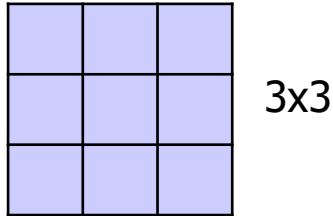
## Upsampling

### Transpose convolution

*Learnable Upsampling, also known as: Upconvolution, or Deconvolution*

Stride: 3

Padding: 1



4x4

Parameter sharing

## Neural Network Segmentation - Basic structure

Convolutions

Downsampling

Upsampling

### Skip connections

*Trade-off between classification and localization*

- High level features from later in the network, enable high classification performance, since they are more discriminative and contain more useful semantic information.
- On the other hand, those deep features have low resolution and, thus pose a problem for localization performance.
- Problem of vanishing gradients was solved with skip connections.

# Neural Network Segmentation - Basic structure

Convolutions

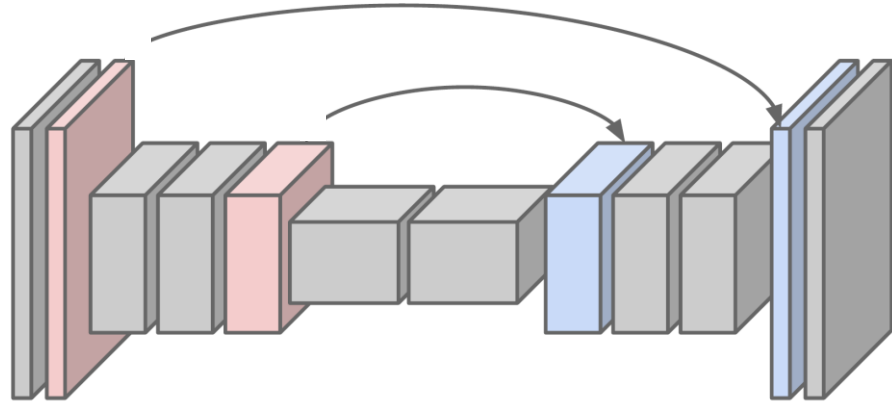
Downsampling

Upsampling

## Skip connections

Combining low-level features,  
which have high localization  
accuracy

With the high-level features,  
which have are descriptive but  
low-resolution.



<http://cs231n.github.io/>

# Neural Network Segmentation - Optimization

## Last layer

Last Layer results in a tensor with  $H \times W$  image resolution and a depth of

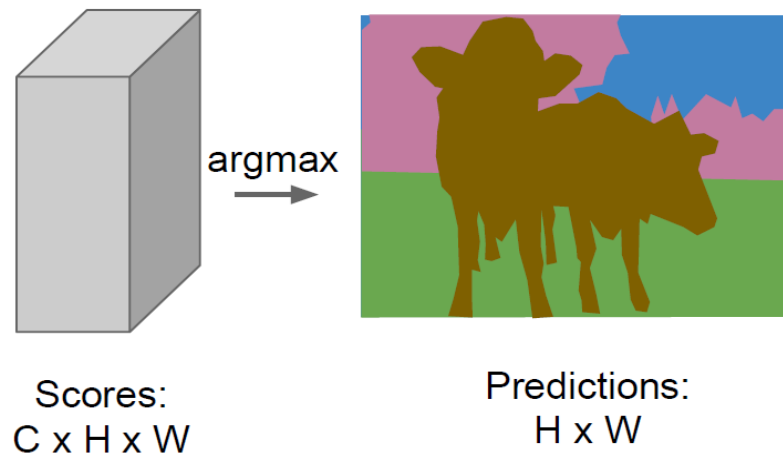
$C$ : Number of classes to segment.

The last layer should encode the values into a range of values of  $(0;1)$ .

Usually either with **softmax** or **sigmoid** function.

Cross-entropy

Dice-coefficient



<http://cs231n.github.io/>

# Neural Network Activation Functions - Review

## Sigmoid

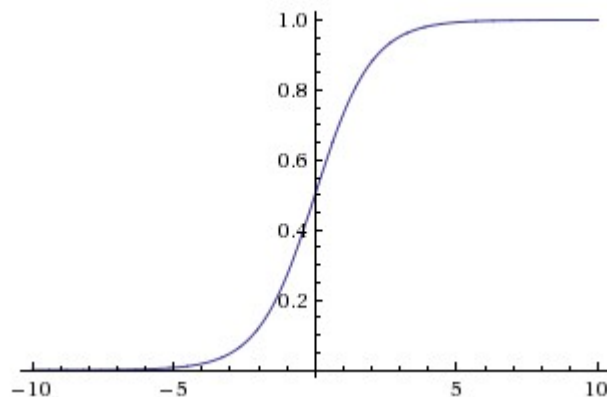
Binary classification only.

The probability for different classes does not need to sum to one.

Simply take the highest class probability.

| class      | model out<br>( $x_i$ ) | prob out<br>( $\hat{y}_i$ ) |
|------------|------------------------|-----------------------------|
| cow        | 2.0                    | 0.88                        |
| grass      | 1.0                    | 0.73                        |
| background | -3.0                   | 0.47                        |

$$\hat{y}_i = 1 - \frac{1}{1 + e^{-x_i}}$$



<http://cs231n.github.io/neural-networks-1/>

# Neural Network Activation Functions - Review

## Softmax

Normalized exponential function

Often used for multi-class segmentation.

Probability sum will be 1.

$$\hat{y}_i = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

| class      | model out<br>( $x_i$ ) | prob out<br>( $\hat{y}_i$ ) |
|------------|------------------------|-----------------------------|
| cow        | 2.0                    | 0.72                        |
| grass      | 1.0                    | 0.27                        |
| background | 0.1                    | 0.01                        |



Last layer

### Binary Cross-entropy

$$L = -\frac{1}{n} \sum_{i=1}^n (\tilde{y}_i \cdot \log(\hat{y}_i) + (1 - \tilde{y}_i) \cdot \log(1 - \hat{y}_i))$$

$\tilde{y}_i$  = *expected pixel class (boolean)*

$\hat{y}_i$  = *predicted pixel class,  $\hat{y}_i \in (0,1)$*

Dice-coefficient

Last layer

### Binary Cross-entropy

$$L = -\frac{1}{n} \sum_{i=1}^n (\tilde{y}_i \cdot \log(\hat{y}_i) + (1 - \tilde{y}_i) \cdot \log(1 - \hat{y}_i))$$

$\tilde{y}_i$  = expected pixel class (boolean)

$\hat{y}_i$  = predicted pixel class,  $\hat{y}_i \in (0,1)$

Dice-coefficient

$\tilde{y}_i \cdot \log(\hat{y}_i)$  = error from positive class

$(1 - \tilde{y}_i) \cdot \log(1 - \hat{y}_i)$  = error from negative class

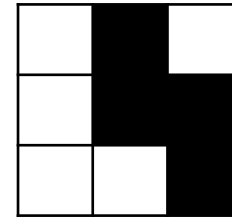
# Neural Network Segmentation - Optimization

Last layer

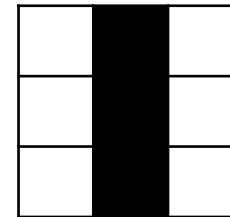
Binary Cross-entropy

$$L = -\frac{1}{n} \sum_{i=1}^n (\tilde{y}_i \cdot \log(\hat{y}_i) + (1 - \tilde{y}_i) \cdot \log(1 - \hat{y}_i))$$

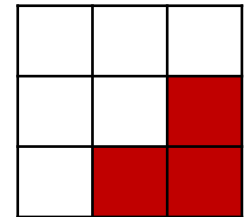
Dice-coefficient



prediction



ground  
truth



error 0.33

# Neural Network Segmentation - Optimization

Last layer

Binary Cross-entropy

## Dice-coefficient

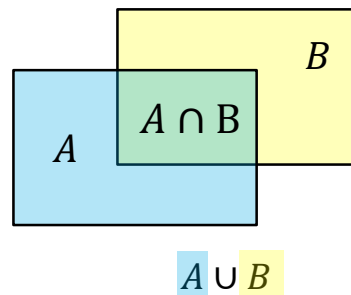
Similar to the IoU (Intersection over union)

More robust with respect to imbalanced classes.

$$L = 1 - \frac{2 \sum_i^n \tilde{y}_i \cdot \hat{y}_i}{\sum_i^n \tilde{y}_i + \sum_i^n \hat{y}_i}$$

$\tilde{y}_i$  = expected pixel class (boolean)

$\hat{y}_i$  = predicted pixel class



# Neural Network Segmentation - Optimization

Last layer

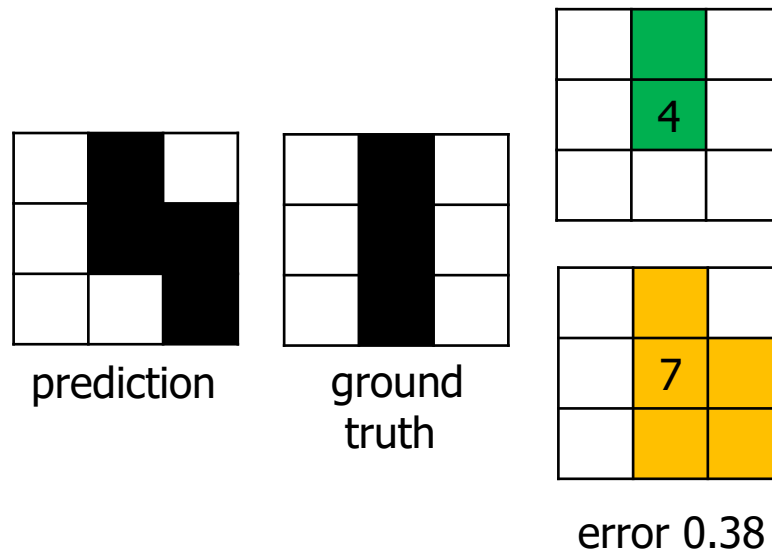
Binary Cross-entropy

**Dice-coefficient**

Similar to the IoU (Intersection over union)

More robust with respect to imbalanced classes.

$$L = 1 - \frac{2 \sum_i^n \tilde{y}_i \cdot \hat{y}_i}{\sum_i^n \tilde{y}_i^2 + \sum_i^n \hat{y}_i^2}$$

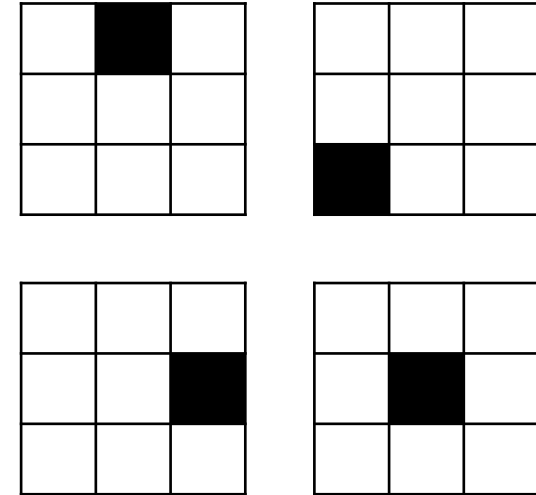


### Thought experiment (at home)

Assumption:

The maximum number of class 1 pixels in a single sample is 1.

This simulates an extreme class imbalance ratio of 1 to 8.



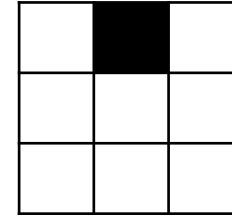
ground  
truth

### Thought experiment (at home)

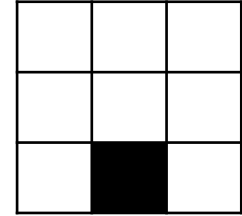
Assumption:

The maximum number of class 1 pixels in a single sample is 1

What is the expected error if the model tries to predict a single pixel of class 1 and fails?



prediction



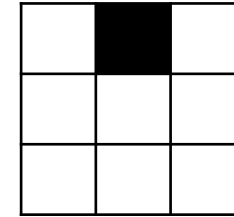
ground  
truth

### Thought experiment (at home)

Assumption:

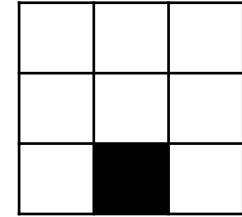
The maximum number of class 1 pixels in a single sample is 1

What is the expected error if the model tries to predict a single pixel of class 1 and fails?



prediction

**BCE**  
**0.22**



ground  
truth

**DL**  
**0.66**



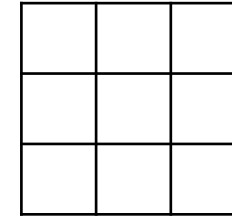
### Thought experiment (at home)

Assumption:

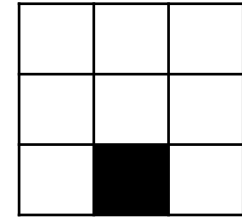
The maximum number of class 1 pixels in a single sample is 1

What is the expected error if the model tries to predict a single class 1?

What is the maximal expected error if the model predicts class 2 only?



prediction



ground  
truth

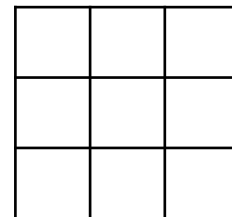
### Thought experiment (at home)

Assumption:

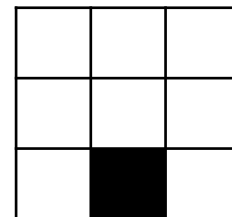
The maximum number of class 1 pixels in a single sample is 1

What is the expected error if the model tries to predict a single class 1?

What is the maximal expected error if the model predicts class 2 only?



prediction



ground  
truth

**BCE**  
**0.11**

**DL**  
**0.5**

### **Thought experiment (at home)**

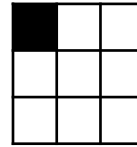
How high is the pressure to get locked in a local minimum if predictions are initially random?

### Thought experiment (at home)

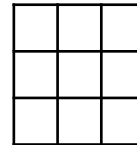
How high is the pressure to get locked in a local minimum if predictions are initially random?

Binary cross-entropy

$$\begin{aligned} 0.22 * 8 + 0.00 * 1 &= 1.76 \\ 0.11 * 9 &= 1.00 \end{aligned}$$



**43 %**

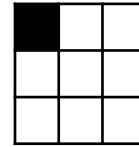


### Thought experiment (at home)

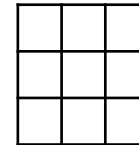
How high is the pressure to get locked in a local minimum if predictions are initially random?

Dice-loss

$$\begin{aligned} 0.66 * 8 + 0.00 * 1 &= 5.28 \\ 0.50 * 9 &= 4.50 \end{aligned}$$



**15 %**



### How to deal with imbalanced classes

Choose **dice-loss** over cross-entropy.

### How to deal with imbalanced classes

Choose dice-loss over cross-entropy.

**Balance** your cross-entropy according to the class imbalance.

*In our case  $\beta = 7/8$*

$$L = -\frac{1}{n} \sum_{i=1}^n (\beta \cdot \tilde{y}_i \cdot \log(\hat{y}_i) + (1 - \beta)(1 - \tilde{y}_i) \cdot \log(1 - \hat{y}_i))$$

## Neural Network Segmentation - Datasets and benchmarking

### PASCAL Visual Object Classes

Pixel-wise segmentation of objects from a number of visual object classes in realistic scenes (i.e. not pre-segmented objects).

#### Annotations

Person, animals, vehicles, indoor.

#### Number of samples

6929 Pixel-wise  
instance level annotations.

#### Metric

Mean Intersection over Union.

Image



Objects



<http://host.robots.ox.ac.uk/pascal/VOC/voc2012/index.html>



# Neural Network Segmentation - Datasets and benchmarking

## Common Objects in Context

COCO-Stuff augments 164K images with pixel-level stuff annotations for semantic segmentation.

### Annotations

91 stuff classes (wall, grass, etc.) and 80 thing classes (person, elephant, etc.), as well as captions.

### Number of samples

164000 dense pixel-level annotations and instance level annotations for things.

### Metric

Mean Intersection over Union.



<https://github.com/nightrome/cocostuff>

## Neural Network Segmentation - Datasets and benchmarking

### Cityscapes

The Cityscapes Dataset focuses on semantic understanding of urban street scenes.

#### Annotations

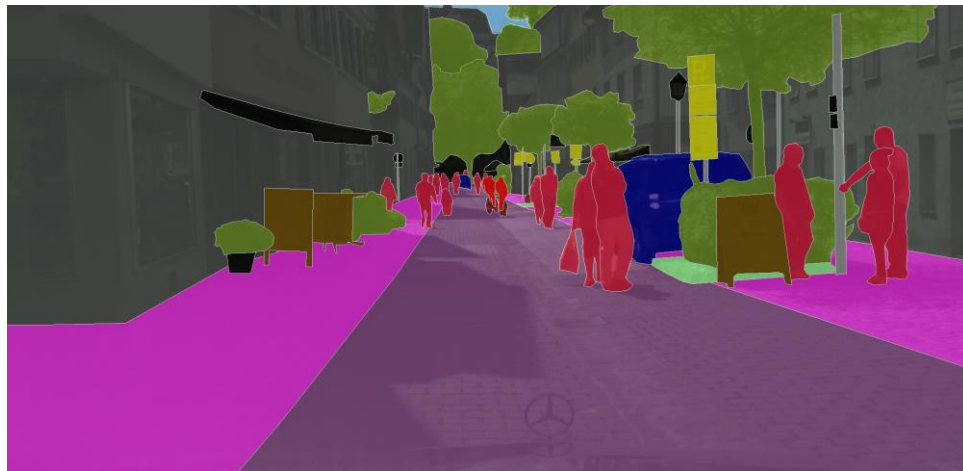
City scene semantic and instance-wise pixel annotations (road, person, pole, etc.).

#### Number of samples

30 classes in 5000 fine and 20000 coarse annotated images.

#### Metric

Mean Intersection over Union  
and Instance Intersection over Union.



<https://www.cityscapes-dataset.com/>

### Architectures over time

|  |      |
|--|------|
| Fully Convolutional Network                | 2015 |
| ParseNet                                   | 2015 |
| Convolutional and Deconvolutional Networks | 2015 |
| U-Net                                      | 2015 |
| Feature Pyramid Network                    | 2016 |
| Mask R-CNN                                 | 2017 |
| DeepLab                                    | 2017 |

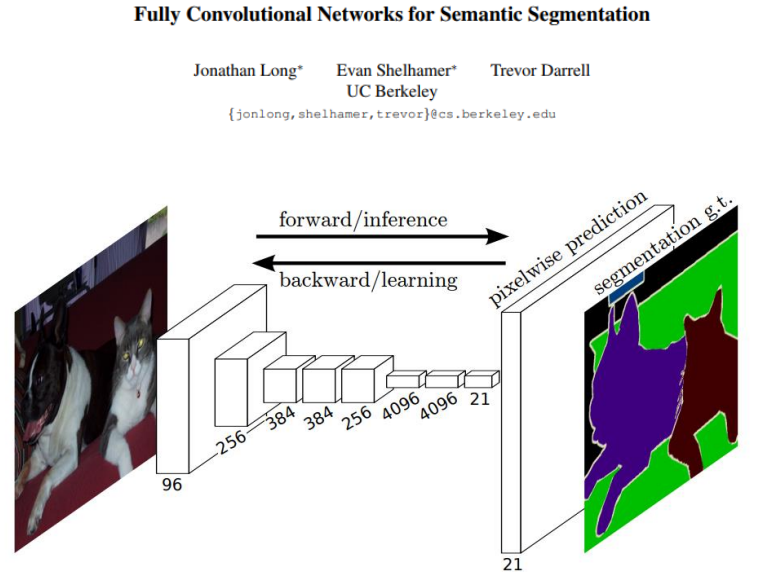
# Neural Network Segmentation - Overview state-of-the-art

## Fully Convolutional Network

First end-to-end trained Fully Convolutional Network for image segmentation.

Transfer Learning approach, modifying well known architectures (such as VGG16).

Ending with an upsampling layer with one channel per class.



[https://people.eecs.berkeley.edu/~jonlong/long\\_shelhamer\\_fcn.pdf](https://people.eecs.berkeley.edu/~jonlong/long_shelhamer_fcn.pdf)

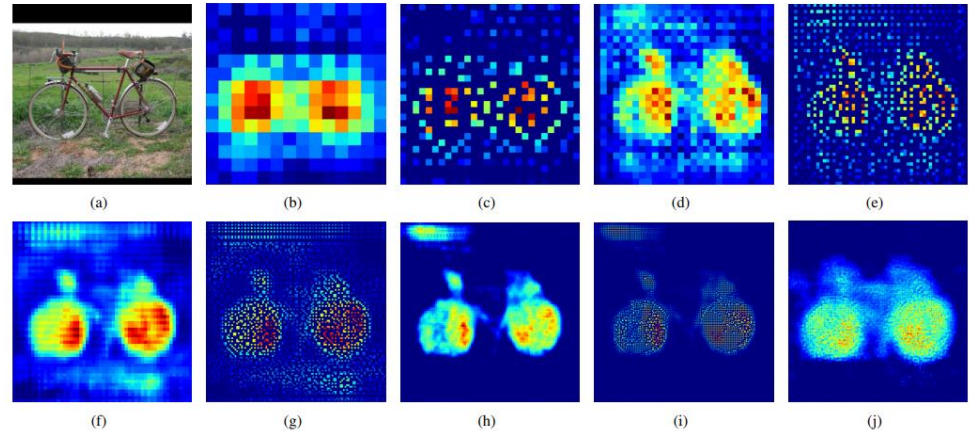
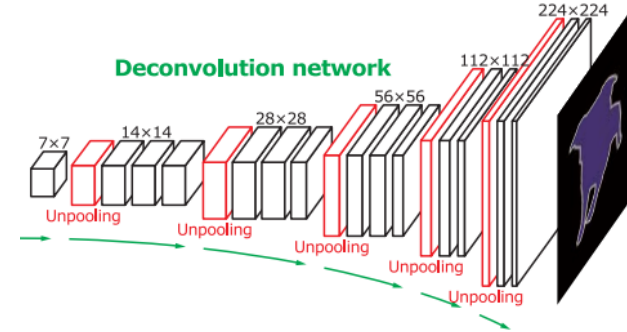
## Neural Network Segmentation - Overview state-of

### Convolutional and Deconvolutional Networks

Introducing an encoder-decoder architecture.

From the convolutional encoding, the deconvolution branch generates a dense pixel-wise class probability map, by successive:

Unpooling, deconvolutions, and rectifications.

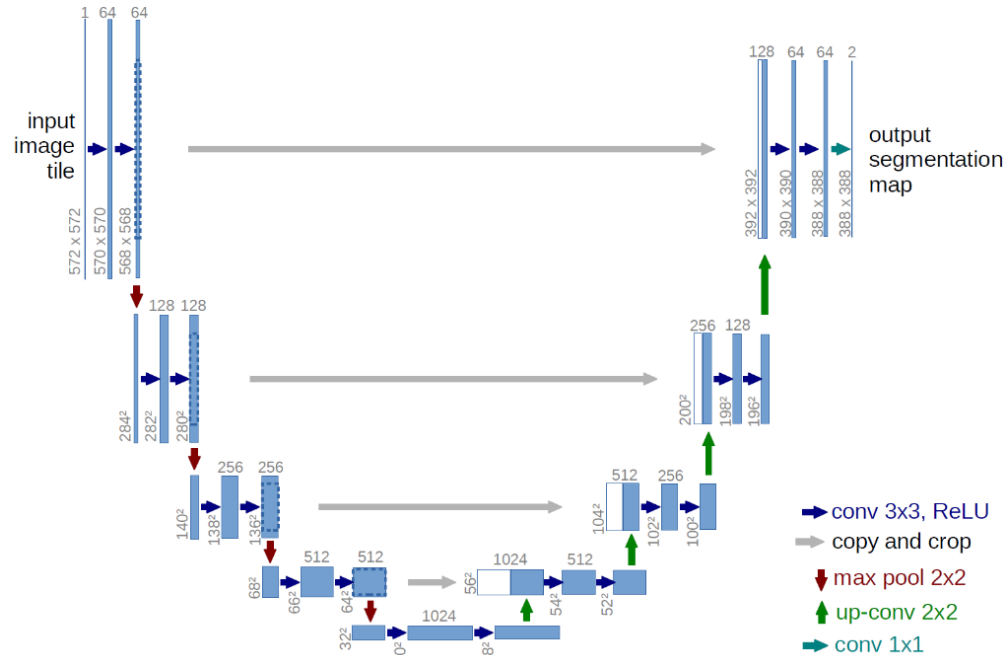


<https://arxiv.org/pdf/1505.04366.pdf>

# Neural Network Segmentation - Deep Dive U-Net

## Deep Dive U-Net

The U-Net is a symmetric, deep convolutional neural network.



<https://lmb.informatik.uni-freiburg.de/people/ronneber/u-net/>

Data  
augmentation

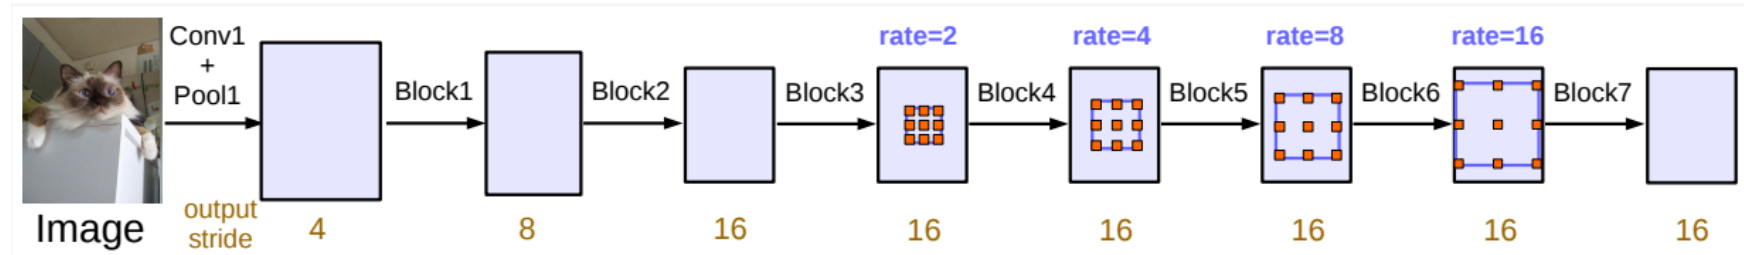
Train  
segmentation  
network

Ventricular  
dimension  
estimation

# Neural Network Segmentation - Overview state-of-the-art

## DeepLabv3

Combining Atrous Convolutions (dilated convolutions) with a pyramidal architecture.



<https://arxiv.org/pdf/1706.05587.pdf>

### Atrous Convolution

Introducing an additional parameter, called the dilation rate.

Defining a spacing between the values in a filter map.

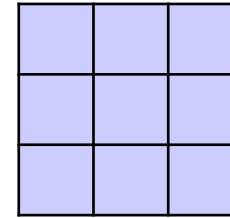


## Neural Network Segmentation - Basic structure

### Atrous Convolution

Introducing an additional parameter, called the dilation rate.

Defining a spacing between the values in a filter map.

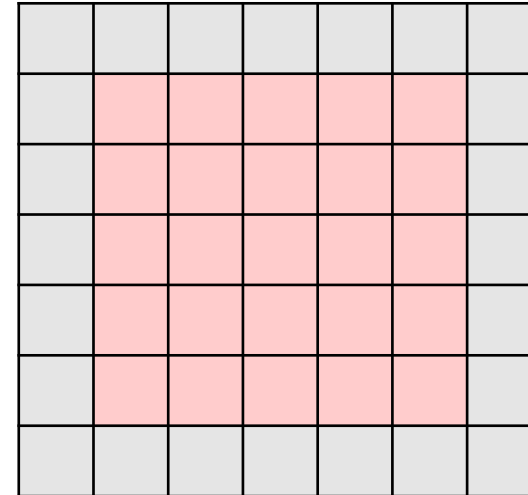


Filter 3x3

**Dilation rate: 2**

Stride: 1

Padding: 1



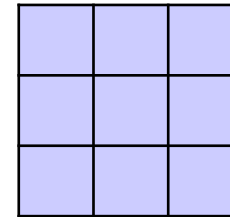
# Neural Network Segmentation - Basic structure

## Atrous Convolution

Introducing an additional parameter, called the dilation rate.

Defining a spacing between the values in a filter map.

This enhances the field of view while keeping the computational cost low.



Filter 3x3

**Dilation rate: 2**

Stride: 1

Padding: 1

